

# Chord

A Scalable Peer-to-Peer  
Lookup Protocol

# About me

- Software Engineer @ MindTouch
- Interests
  - Distributed Systems
  - Algorithms
- Hobbies
  - Martial Arts
  - Snowboarding
  - Diving
  - Mountain Biking




# What is Chord?

- Distributed Hash Table
- Hash Table
  - Mapping of Key -> Value
  - Dictionary
  - Map
  - HashMap

# Examples

- DNS
- K/V Stores (Memcache, Redis, DynamoDB)
- GlusterFS
- AWS S3

WHICH KEY GOES TO WHICH NODE?



# Prior Art

- Central Repository (Napster)
- Query Flooding (Gnutella)
- Hierarchy (Kazaa, modern Gnutella)

# How to organize Nodes

- Up to 90's
  - Master-Slave model (Napster)
  - Unequal responsibilities
- More recently
  - P2P (Gnutella)
  - Equal responsibilities

# Costs

	Memory	Lookup Latency	# Messages
Napster	$O(N)$	$O(1)$	$O(1)$
Gnutella	$O(N)$	$O(N)$	$O(N)$



# Chord Design Goals

- Good Performance
- Simple and correct protocol
- Degrades gracefully under failure
- Scalable to large number of nodes
- No naming structure for keys
- Load balanced
- Decentralized (P2P)
- Available

# Why?

- Distributed Storage System
  - Keys: filename
  - Values: node responsible for storing the file
  - Cooperative mirroring
- Distributed Indexes
  - Keys: search terms
  - Values: nodes containing files with those terms
- Large-Scale Combinatorial Search
  - Keys: candidate solution
  - Values: machines responsible for solution

What data structure is optimized for lookup speeds?

# Hashing

- Maps  $n$ -bit key into  $k$ -buckets
- Function  $\text{fn}(\text{key}) = \text{node identifier}$
- Goals
  - Low cost
  - Deterministic
  - Load balanced
- Simplest Hash Function (I can think of)
  - $\text{fn}(\text{key}) = \text{key} \% \text{\#nodes}$

# Modulo Hashing


- Good Performance
- Simple and Correct (provably) Protocol
- ~~Degrades gracefully under failure~~
- Scalable to large number of nodes
- No naming structure for keys
- ~~Load balanced~~
- Decentralized (P2P)
- ~~Available~~



# Consistent Hashing


# Consistent Hashing


- Organize Nodes in a Ring (Structured P2P)
- Maintains neighbors
- Lookups follow neighbor links clockwise around the ring
- Nodes and Keys are “hashed” onto the Ring





- Maintains a doubly linked list (successor/predecessor)
- Separate Join/Leave protocol





- Even distribution
- Little disruption as nodes join/leave

# Lookup

*// ask node  $n$  to find the successor of  $id$*

**$n$ .find\_successor( $id$ )**

**if ( $id \in (n, \text{successor}]$ )**

**return successor;**

**else**

~~$n' = \text{closest\_preceding\_node}(id);$~~


**return  $n'$ .find\_successor( $id$ );**

$n$ .get\_successor()



# Lookup Complexity?



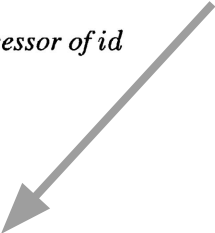


# Lookup

```
// search the local table for the highest predecessor of id  
n.closest_preceding_node(id)  
  for  $i = m$  downto 1  
    if ( $\text{finger}[i] \in (n, id)$ )  
      return  $\text{finger}[i]$ ;  
return  $n$ ;
```

*// ask node n to find the successor of id*

```
n.find_successor(id)  
  if ( $id \in (n, \text{successor}]$ )  
    return successor;  
  else  
     $n' = \text{closest\_preceding\_node}(id)$ ;  
    return  $n'.\text{find\_successor}(id)$ ;
```




# Lookup Complexity?




# Join/Leaves


# Join




(a)



(b)



(c)



(d)

*// join a Chord ring containing node  $n'$ .*

**$n.join(n')$**

*predecessor = nil;*

*successor =  $n'.find\_successor(n)$ ;*

*// called periodically. verifies  $n$ 's immediate*

*// successor, and tells the successor about  $n$ .*

**$n.stabilize()$**

*$x = successor.predecessor$ ;*

**if** ( $x \in (n, successor)$ )

*successor =  $x$ ;*

*successor.notify( $n$ );*

*//  $n'$  thinks it might be our predecessor.*

**$n.notify(n')$**

**if** (*predecessor is nil or  $n' \in (predecessor, n)$* )

*predecessor =  $n'$ ;*

*// called periodically. refreshes finger table entries.*

*// next stores the index of the next finger to fix.*

***n.fix\_fingers()***

*next = next + 1;*

***if (next > m)***

*next = 1;*

*finger[next] = find\_successor( $n + 2^{next-1}$ );*

# Fingers can be out of date!

- Lookups are still correct
- Degrades to  $O(N)$  in the worst case!

# How do JOIN's affect operations


- Lookups are correct but may be slow (due to outdated pointers)
- Some lookups may fail
  - Data is in transit
  - Incorrect successor pointers
- SOLUTION:
  - Higher level software should handle it

# Leaving the Ring!

**TWO WEEKS  
NOTICE**







# Leaving the Ring (w/ proper 2 weeks notice)

- Notify your predecessor/successor by sending them your successor/predecessor
- Transfer them your data ahead of time (if you are nice)

# Failure handling

- The last hop is always from predecessor  $\rightarrow$  successor
- It is the predecessor's job to notice if the successor has failed
  - But what next?
- Instead of keeping a single successor, keep a list of successors
- When you notice your successor has failed, notify successor+1, and cut your successor out of the circle

## Being accepting of new predecessors

*// called periodically. checks whether predecessor has failed.*

***n.check\_predecessor()***

***if (predecessor has failed)***


***predecessor = nil;***

# Chord does not...

- Prevent split brain
- Prevent rings that loop multiple times around
- SOLUTION:
  - Sample the ring topology every so often and fix it (thanks?)
  - Though they do hypothesize, an ordinary series of JOINS/LEAVES do not cause this.

# Analysis

(a)



- $10^4$  nodes
- $5 * 10^5$  keys

TABLE II  
 PATH LENGTH AND NUMBER OF TIMEOUTS EXPERIENCED BY A LOOKUP AS  
 FUNCTION OF THE FRACTION OF NODES THAT FAIL SIMULTANEOUSLY. THE  
 FIRST AND 99TH PERCENTILES ARE IN PARENTHESES. INITIALLY, THE  
 NETWORK HAS 1000 NODES

<b>Fraction of failed nodes</b>	<b>Mean path length (1st, 99th percentiles)</b>	<b>Mean num. of timeouts (1st, 99th percentiles)</b>
0	3.84 (2, 5)	0.0 (0, 0)
0.1	4.03 (2, 6)	0.60 (0, 2)
0.2	4.22 (2, 6)	1.17 (0, 3)
0.3	4.44 (2, 6)	2.02 (0, 5)
0.4	4.69 (2, 7)	3.23 (0, 8)
0.5	5.09 (3, 8)	5.10 (0, 11)



TABLE III

PATH LENGTH AND NUMBER OF TIMEOUTS EXPERIENCED BY A LOOKUP AS FUNCTION OF NODE JOIN AND LEAVE RATES. FIRST AND 99TH PERCENTILES ARE IN PARENTHESES. THE NETWORK HAS ROUGHLY 1000 NODES

<b>Node join/leave rate (per second/per stab. period)</b>	<b>Mean path length (1st, 99th percentiles)</b>	<b>Mean num. of timeouts (1st, 99th percentiles)</b>	<b>Lookup failures (per 10,000 lookups)</b>
0.05 / 1.5	3.90 (1, 9)	0.05 (0, 2)	0
0.10 / 3	3.83 (1, 9)	0.11 (0, 2)	0
0.15 / 4.5	3.84 (1, 9)	0.16 (0, 2)	2
0.20 / 6	3.81 (1, 9)	0.23 (0, 3)	5
0.25 / 7.5	3.83 (1, 9)	0.30 (0, 3)	6
0.30 / 9	3.91 (1, 9)	0.34 (0, 4)	8
0.35 / 10.5	3.94 (1, 10)	0.42 (0, 4)	16
0.40 / 12	4.06 (1, 10)	0.46 (0, 5)	15

# Credits & Resources

- Chord: [https://pdos.csail.mit.edu/papers/chord:sigcomm01/chord\\_sigcomm.pdf](https://pdos.csail.mit.edu/papers/chord:sigcomm01/chord_sigcomm.pdf)
- Great slides on Chord: <http://slideplayer.com/slide/10698351/>
- Great list of papers: <http://cseweb.ucsd.edu/classes/sp14/cse223B-a/syllabus.html>